

DENTFLUENCE · V1 · VERSION 4 — 28 AUGUST

Decision Log

Every decision on this round is now closed. Eleven answers go into work as written; N2 is dropped from V1 at your word and moves to V1.1. Nothing in this document is waiting on you any more — the one live item left is C1, and that is paperwork with Meta rather than a choice. What follows is the record of what was decided, what it costs, and the three things that stay in scope regardless.

11 SETTLED	1 DROPPED TO V1.1	1 LAUNCH- CRITICAL RISK	~10 WORKING DAYS LEFT
0 DAYS OF SLACK			

PART 1 · CLOSED — GOING INTO WORK AS WRITTEN

Nothing below needs another word. C2, C3, C4, C5, H1, H2, H3 and H4 were closed on 28 August; N3 and N1 close here; N2 closes in Part 2.

N3

Park Smart Presentation too?

CLOSED — AND IT COSTS NOTHING

YOU SAID

"Yes, don't touch treatment plan presentation now."

Taken together with C3, that gives one clean position: **the Treatment Plan tab's present-and-print workflow is untouched, and neither presentation engine gets developed further.**

WHAT I RECOMMEND NOW, HAVING THOUGHT ABOUT IT AGAINST THE FREEZE

When I first suggested parking Smart Presentation I hadn't weighed it against your own zero-slack directive. Doing so changes my answer: **the cheapest safe action is no action.** Smart Presentation has no menu item and nothing in the Treatment Plan tab links to it — no member of staff can reach it. Removing it is itself a code change under a freeze, for a screen nobody can open.

So: leave it alone, record it as V1.1 cleanup. The same logic applies to C3 — before anyone edits anything, **check what `case_acceptance.enabled` is set to in production.** It is true in the local copy but production was never verified. If it is already false there, C3 is closed with **zero work** and the Case Journey button is already absent from your live plan cards.

NET EFFECT ON THE PLAN

C3 and N3 together: **somewhere between zero and fifteen minutes of work,** depending on one production setting. Down from the half-day I had them at.

N1

"V1 should show all important KPIs needed for dental business"

CLOSED — YOU ALREADY ANSWERED THIS

YOU SAID

"V1 should show all important KPI needed for dental business."

You settled this on **26 August**, and more precisely than a new list would. Your directive froze **28 KPIs**, "**no additions**", and put the important ones in Tier 1:

TIER 1 — OWNER / FINANCIAL (8)	ALREADY IN YOUR P0 QUEUE AS
Collections · Outstanding · Expenses · GST / taxable sales	Served by the money engine + <code>ReportMetricsService</code>
Revenue by category	G-01 — next after G-33
Doctor-wise revenue	G-02
Cash / bank	G-04
Profit	G-06

So "all important KPIs" is not new work — **it is your existing P0 queue**. Four of the eight Tier-1 KPIs are literally the next four money items you already scheduled. Designing a fresh KPI set now would breach your own rule against "new KPI concepts", and would fork the definitions your principle 4 exists to prevent.

WHAT I'LL DO INSTEAD OF A NEW BUILD

When G-01, G-02, G-04 and G-06 land, the Analytics overview surfaces the **eight Tier-1 KPIs**, every one of them read through `ReportMetricsService` — your principle 2, one source of truth, no competing calculation. That is a display job sitting on top of work already in the queue, not a parallel project.

REVISED COST

About a day of display work after G-06 lands — not the 3–5 days I quoted when I thought you wanted a new KPI set. And it removes the duplicate Analytics door at the same time (C5).

PART 2 · N2 — DROPPED, AND WHAT THAT DOES NOT REMOVE

You read the arithmetic and made the call. That closes the last open decision on this round: **nothing in Part 1 or Part 2 is waiting on you any more.** The only live item left in this document is C1, and that one is paperwork with Meta, not a decision.

N2

Projected vs actual — dropped from V1

CLOSED — YOU SKIPPED IT

YOU SAID

"If N2 is too much issue, skip it."

Skipped. Projected-versus-actual — yesterday, week and month — is **out of V1 entirely** and moves to the V1.1 list. No lever gets pulled, no P0 gets dropped, the date does not move, and the queue goes back to exactly what your 26 August directive said it was.

This is the right instinct and I want to name why, because you will be tempted by the same shape of request again before 14 September: N2 was a *reporting* improvement bought with days that belong to *verifying money*. Your own principle — a wrong KPI is worse than a missing one — makes that trade a bad one at this point in the calendar. The tracking loop is worth having. It is not worth having at the cost of shipping G-04 or G-06 unverified.

WHAT THIS DOES NOT REMOVE — THREE THINGS STAY EXACTLY WHERE THEY WERE

1 · H4 stays. The manager still fills *amount to collect* in the morning huddle. That field is already built and already saves; dropping N2 removes the scoreboard, not the habit. Start the habit on day one anyway — see the Rollout Plan — because the day you do build the comparison, it can only read history that was actually entered.

2 · H3 stays, and it was never part of this dispute. The Huddle report reading the wrong money source is a *defect*, not a feature request. It is already on your register as **G-32**, sized 1–2 days, and it is unaffected by skipping N2. Weekly and monthly still need to show the same figures as daily.

3 · The numbers are not lost. Every projection the manager enters is stored on the appointment row from the first morning you start. Whenever you build the comparison — V1.1, October, whenever — it reads that history. Nothing is being thrown away by waiting; you are only deferring the display.

STANDING OFFER, NO DECISION NEEDED TODAY

The half-day version is still on the shelf: one line on the Huddle board reading "Yesterday — projected ₹18,400 · collected ₹16,900 · from 9 of 11 appointments." If the P0 queue lands ahead of schedule — or if you find after two weeks of huddles that the manager's estimates are drifting and you want them scored in

front of the team — say the word and it goes in. Until then it sits in V1.1 with the rest of N2, and I will not raise it again.

THE BUDGET AFTER THIS DECISION

P0 work committed · 11.5–13 days

■ Committed P0 queue (G-33 → G-01 → G-02 → G-07 → G-04 → G-05 → G-06 → G-08)

Back to the plan as written on 26 August — which is still **~11.5–13 days of work against roughly 10 working days**. Skipping N2 stopped the overrun growing; it did not create slack. The queue is still at or past capacity, the code-freeze gate has not lifted, and G-01 has not started. If something has to give later, your directive already names the order: **G-09 first, then G-13** — and that decision is better made in the first week of September than in the last.

PART 3 · C1 — THE WHATSAPP CHANNEL

You've moved this from "a dependency" to "must be bulletproof before launch, and it's our differentiator". That changes its status: it is now the **single most likely thing to break 14 September**, and not for a reason any amount of coding fixes. Three things you need to know before you plan around it.

C1a

Meta's clock is not your clock — and it hasn't started

START TODAY

YOU SAID

"Need this channel full proof before launch, but expecting to rise by 14th September. This is going to be a differentiating feature."

Getting a WhatsApp Business number live is a paperwork chain with external waiting built into every link, and today the clinic holds none of it — no phone number ID, no business account ID, no token. Each step below can bounce back for a document mismatch or a template wording rejection, and each bounce restarts that step's clock.

STEP	WHO CONTROLS THE CLOCK	TYPICAL WAIT
Meta Business verification (documents, business name match)	Meta	Days to a few weeks — the long pole, and the one that bounces
WhatsApp Business Account + sending number verification	Meta	Hours to days
7 message templates submitted and approved	Meta	24–48h each, rejections common on wording
Credentials into production config	You	Under an hour
Controlled first send	You	One supervised day

There are 17 calendar days to 14 September. If verification runs long — and it is the step most likely to — no amount of development recovers it.

THE ONE THING TO DO TODAY

Start Meta Business verification **now**, before any of the code decisions in this document. It is not development work — it is documents, and it is the only item on

the whole V1 plan where waiting is someone else's decision. Everything else in this log can be done in the time the verification is sitting in a queue.

c1b

"Full proof" can be achieved before Meta answers — that's the good news

DO THIS IN PARALLEL

Practice mode is not a limitation here, it is the test rig. Every message is composed in full, run through the consent gate, and written to the patient's record — it is only the final handover to Meta that is skipped. **So you can prove correctness completely without credentials.**

The proving tools already exist as commands: `whatsapp:test` for a plain message and `whatsapp:template` for a template, both described in the code as dry-run safe.

WHAT "FULL PROOF" SHOULD MEAN — PROPOSED, SO IT'S A BAR YOU CAN ACTUALLY CLEAR

- **Right person.** Every send resolves to the intended patient and the correctly formatted number, including numbers stored without a country code.
- **Right words.** All seven templates render with real patient data — name, date, time, amount — with no blank or wrong variable.
- **Right permission.** A patient who has not granted WhatsApp consent is refused. A patient who granted service consent but not marketing is refused the promotional template and allowed the reminder.
- **No duplicates.** Running the same job twice does not queue the same message twice.
- **Right record.** Every attempt appears on the patient's timeline whether it sent, was refused by consent, or failed.
- **Fails safely.** With bad credentials the app logs and moves on — it does not crash a scheduled job or block reception.

Six checks, no Meta account required. Do them next week and the day the credentials arrive is a configuration change, not a gamble.

C1c

The automation behind the differentiator is thinner than it looks

READ BEFORE YOU POSITION IT

I checked what actually sends on its own, because "differentiating feature" is a promise to patients and to yourself. Seven templates are configured. **Two of them are wired to send automatically.**

TEMPLATE	SENDS AUTOMATICALLY?	WHAT ACTUALLY HAPPENS TODAY
Appointment reminder	Yes	Scheduled job at 10:00 daily
Review request	Yes	Scheduled job at 11:00 daily, after completed visits
Recall due	No	The recall engine builds a <i>call list</i> for reception. It stores the patient's WhatsApp number but sends nothing — a person presses the button, one patient at a time
Payment reminder	No	Template exists; nothing sends it
Lab ready	No	Template exists; nothing sends it
Appointment confirmation	No	Template exists; nothing sends it
Festive offer	No	Manual, and marketing-consent gated

WHY THIS MATTERS TO HOW YOU SELL IT

Recall is the one that brings patients back — and it is manual. On 14 September, switching WhatsApp on gives you automatic appointment reminders and automatic review requests. Everything else is a button a human presses.

That is still a real differentiator against a clinic doing none of it, and manual-with-one-click is a perfectly good V1. But if the pitch is "the system chases your patients for you", the honest version today is "the system tells your reception exactly who to chase, and gives them one button to do it".

Automating recall sends is a V1.1 conversation, not a V1 one — and it should not be started inside this freeze.

CONFIRM YOU'RE HAPPY LAUNCHING WITH REMINDERS + REVIEWS
AUTOMATED

WHERE THE WORK STANDS AFTER THIS ROUND

JOB	FROM	SIZE	STATUS
Start Meta Business verification	C1	Paperwork, not dev	Today — the only true blocker
Remove the fake Marketing column	H1	½ day	Ready
Hide "View All Yesterday"	H2	1 hour	Ready
Lock clinical files + DPDP screens	C2, C4	½ day together	Ready
Check <code>case_acceptance.enabled</code> in production	C3, N3	Minutes	Ready — may be zero work
Remove duplicate Finance → Analytics door	C5	½ day	Ready
Fix the Huddle report's money source	H3	1–2 days	Ready (already G-32)
Prove the WhatsApp channel in practice mode	C1b	1 day	Ready — no Meta needed
Yesterday projected vs actual on the board	N2	—	Dropped — moved to V1.1
Tier-1 KPIs on the Analytics overview	N1	~1 day	After G-06 — inside existing queue

JOB	FROM	SIZE	STATUS
Unblocked and ready now		~3.5 days	Fits before the P0 queue resumes

None of it has been started. The code-freeze gate from 26 August is still down — it lifts when you confirm production is green, and that verification is still step one of everything.

V1 Decision Log · version 4, 28 August 2026. Every claim checked against the running code or against your own 26 August directive before it was written here — including the place where I pushed back on your answer rather than recording it — a push-back you then settled by dropping the item.

Companion documents: the Owner's Handbook, the Marathi Staff Guide, and the Daily Huddle Rollout Plan. Appointments is the next chapter.